

```
In [2]: import pandas as pd  
import numpy as np  
from PIL import Image
```

```
In [5]: img = Image.open('data-cleansing.jpg')    # Open image as PIL image object  
img
```

Out[5]:



ABOUT THE DATASET

The dataset was downloaded from a random github account. The aim of this project is to use various methods available through python Pandas to clean up the data to a state where it is fit for Analysis and Visualization. The visualization aspect will be a separate project. This project is dedicated solely to Data Cleaning

READING HOUSING CSV DATASET INTO JUPYTER

```
In [126]: df = pd.read_csv('nash_housing_data.csv')
```

DISPLAYING FIRST FIVE ROWS

In [127]:

df.head(5)

Out[127]:

	UniqueID	ParcelID	LandUse	PropertyAddress	SaleDate	SalePrice	LegalReference	SoldAsVa
0	2045	007 00 0 125.00	SINGLE FAMILY	1808 FOX CHASE DR, GOODLETTSVILLE	09-Apr- 13	240000	20130412- 0036474	
1	16918	007 00 0 130.00	SINGLE FAMILY	1832 FOX CHASE DR, GOODLETTSVILLE	10-Jun- 14	366000	20140619- 0053768	
2	54582	007 00 0 138.00	SINGLE FAMILY	1864 FOX CHASE DR, GOODLETTSVILLE	26-Sep- 16	435000	20160927- 0101718	
3	43070	007 00 0 143.00	SINGLE FAMILY	1853 FOX CHASE DR, GOODLETTSVILLE	29-Jan- 16	255000	20160129- 0008913	
4	22714	007 00 0 149.00	SINGLE FAMILY	1829 FOX CHASE DR, GOODLETTSVILLE	10-Oct- 14	278000	20141015- 0095255	

UNDERSTANDING FEATURES OF THE DATASET

In [128]:

df.shape

Out[128]: (56477, 19)

In [129]:

df.describe()

Out[129]:

	UniqueID	Acreage	LandValue	BuildingValue	TotalValue	YearBuilt	Be
count	56477.000000	26015.000000	2.601500e+04	2.601500e+04	2.601500e+04	24163.000000	24157
mean	28334.001133	0.498923	6.906856e+04	1.607847e+05	2.323754e+05	1963.744899	3
std	16352.590651	1.570454	1.060401e+05	2.067999e+05	2.810643e+05	26.542982	(
min	0.000000	0.010000	1.000000e+02	0.000000e+00	1.000000e+02	1799.000000	(
25%	14186.000000	0.180000	2.100000e+04	7.590000e+04	1.028000e+05	1948.000000	3
50%	28313.000000	0.270000	2.880000e+04	1.114000e+05	1.485000e+05	1960.000000	3
75%	42513.000000	0.450000	6.000000e+04	1.807000e+05	2.683500e+05	1983.000000	3
max	56635.000000	160.060000	2.772000e+06	1.297180e+07	1.394040e+07	2017.000000	1'

In [130]:

From the counts we can see that some columns have empty cells

In [131]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 56477 entries, 0 to 56476
Data columns (total 19 columns):
#   Column                Non-Null Count  Dtype
---  -
0   UniqueID              56477 non-null  int64
1   ParcelID              56477 non-null  object
2   LandUse               56477 non-null  object
3   PropertyAddress       56448 non-null  object
4   SaleDate              56477 non-null  object
5   SalePrice             56477 non-null  object
6   LegalReference        56477 non-null  object
7   SoldAsVacant          56477 non-null  object
8   OwnerName             25261 non-null  object
9   OwnerAddress          26015 non-null  object
10  Acreage                26015 non-null  float64
11  TaxDistrict           26015 non-null  object
12  LandValue             26015 non-null  float64
13  BuildingValue         26015 non-null  float64
14  TotalValue            26015 non-null  float64
15  YearBuilt             24163 non-null  float64
16  Bedrooms              24157 non-null  float64
17  FullBath              24275 non-null  float64
18  HalfBath              24144 non-null  float64
dtypes: float64(8), int64(1), object(10)
memory usage: 8.2+ MB
```

In [132]: `df.isnull().sum()`

```
Out[132]: UniqueID              0
ParcelID              0
LandUse              0
PropertyAddress       29
SaleDate              0
SalePrice             0
LegalReference        0
SoldAsVacant          0
OwnerName            31216
OwnerAddress          30462
Acreage              30462
TaxDistrict          30462
LandValue            30462
BuildingValue        30462
TotalValue           30462
YearBuilt            32314
Bedrooms             32320
FullBath             32202
HalfBath             32333
dtype: int64
```

In [133]: *# We can delete rows with null values with this code <df_drop=df.dropna()> but for all nulls in place. I prefer leaving nulls in the dataset and omitting them where the nulls will affect the outcome of the results*

DROPPING DUPLICATE ROWS IF THEY EXIST

```
In [134]: # Let's drop duplicate rows
df=df.drop_duplicates(keep='last')
```

CHECKING IF THERE WERE DUPLICATES We do this by re-checking the shape

```
In [135]: df.shape
```

```
Out[135]: (56477, 19)
```

```
In [136]: # The shape is still the same so there are no duplicates
```

CHANGING DATE SaleDate COLUMN IN TO DATETIME SO PANDAS RECOGNIZES IT AS A DATE COLUMN

```
In [137]: df['SaleDate']=pd.to_datetime(df['SaleDate'])
```

CHECKING

```
In [138]: df.head(1)
```

```
Out[138]:
```

	UniqueID	ParcelID	LandUse	PropertyAddress	SaleDate	SalePrice	LegalReference	SoldAsVa
0	2045	007 00 0 125.00	SINGLE FAMILY	1808 FOX CHASE DR, GOODLETTSVILLE	2013-04- 09	240000	20130412- 0036474	

```
In [139]: #Splitting up the SaleDate colmun by the comma in the string
new=df["PropertyAddress"].str.split(",",n=1,expand=True)
```

ASSIGNING NAMES TO SPLIT PARTS

```
In [140]: df["Street Name"]=new[0]
```

```
In [141]: df["City"]=new[1]
```

NOW LET'S DROPP THE SaleDate COLUMN

```
In [142]: df.drop(columns=["PropertyAddress"],inplace=True)
```

CHECKING NEW COLUMNS

In [143]: `df.columns`

Out[143]: Index(['UniqueID ', 'ParcelID', 'LandUse', 'SaleDate', 'SalePrice', 'LegalReference', 'SoldAsVacant', 'OwnerName', 'OwnerAddress', 'Acreage', 'TaxDistrict', 'LandValue', 'BuildingValue', 'TotalValue', 'YearBuilt', 'Bedrooms', 'FullBath', 'HalfBath', 'Street Name', 'City'], dtype='object')

PRINTING SOME ROWS TO CHECK IF THE COMMA THAT WAS IN THE ADDRESS HAS BEEN REMOVED

In [144]: `df.head(5)`

Out[144]:

	UniqueID	ParcelID	LandUse	SaleDate	SalePrice	LegalReference	SoldAsVacant	OwnerName
0	2045	007 00 0 125.00	SINGLE FAMILY	2013-04- 09	240000	20130412- 0036474	No	FRAZIEF CYRENTHE LYNETT
1	16918	007 00 0 130.00	SINGLE FAMILY	2014-06- 10	366000	20140619- 0053768	No	BONEF CHARLES LESLI
2	54582	007 00 0 138.00	SINGLE FAMILY	2016-09- 26	435000	20160927- 0101718	No	WILSON JAMES E. JOANN
3	43070	007 00 0 143.00	SINGLE FAMILY	2016-01- 29	255000	20160129- 0008913	No	BAKER, JAY R & SUSAN E
4	22714	007 00 0 149.00	SINGLE FAMILY	2014-10- 10	278000	20141015- 0095255	No	POST CHRISTOPHE M. SAMANTHA C

In [145]: `# It has been removed`

LET'S USE SAME METHOD TO SPLIT UP THE OwnerAddress

In [146]: `new1=df["OwnerAddress"].str.split(",",n=1,expand=True)`

In [147]: `df["Owner_Street Name"]=new1[0]`

In [148]: `df["Owner_City"]=new1[1]`

In [149]: `df["Owner_State"]=new1[2]`

```

KeyError                                Traceback (most recent call last)
<ipython-input-149-63183eb32529> in <module>
----> 1 df["Owner_State"]=new1[2]

~\anaconda3\lib\site-packages\pandas\core\frame.py in __getitem__(self, key)
    2900         if self.columns.nlevels > 1:
    2901             return self._getitem_multilevel(key)
-> 2902         indexer = self.columns.get_loc(key)
    2903         if is_integer(indexer):
    2904             indexer = [indexer]

~\anaconda3\lib\site-packages\pandas\core\indexes\range.py in get_loc(self, key, method, tolerance)
    355         return self._range.index(new_key)
    356     except ValueError as err:
--> 357         raise KeyError(key) from err
    358     raise KeyError(key)
    359     return super().get_loc(key, method=method, tolerance=tolerance)

```

In [150]: *# I am getting an error because I have to split the second part. The syntax I use*

In [151]: `df.head(5)`

1	16918	007 00 0 130.00	SINGLE FAMILY	2014-06- 10	366000	20140619- 0053768	No	BO CHARL LE	
2	54582	007 00 0 138.00	SINGLE FAMILY	2016-09- 26	435000	20160927- 0101718	No	WIL JAMES JOA	
3	43070	007 00 0 143.00	SINGLE FAMILY	2016-01- 29	255000	20160129- 0008913	No	BAKER, J. & SUS	
4	22714	007 00 0 149.00	SINGLE FAMILY	2014-10- 10	278000	20141015- 0095255	No	P CHRISTOP SAMANTH	

SPLITTING Owner_City Column

In [152]: `new3=df["Owner_City"].str.split(",",n=1,expand=True)`

In [153]: `df["City"]=new1[0]`

```
In [154]: df["State"]=new1[1]
```

```
In [155]: # Now Let's drop Owner_City column and OwnerAddress
```

```
In [156]: df.drop(columns=["Owner_City"],inplace=True)
```

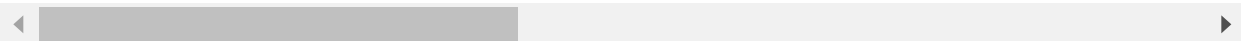
```
In [157]: df.drop(columns=["OwnerAddress"],inplace=True)
```

```
In [158]: df.head(2)
```

Out[158]:

	UniqueID	ParcelID	LandUse	SaleDate	SalePrice	LegalReference	SoldAsVacant	OwnerName
0	2045	007 00 0 125.00	SINGLE FAMILY	2013-04- 09	240000	20130412- 0036474	No	FRAZIER, CYRENTHA LYNETTE
1	16918	007 00 0 130.00	SINGLE FAMILY	2014-06- 10	366000	20140619- 0053768	No	BONER, CHARLES & LESLIE

2 rows × 21 columns



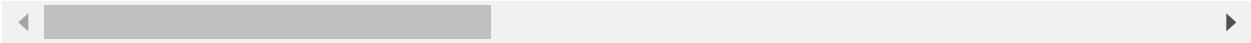
```
In [159]: # Lets rename City and State columns to 'Owner_City' and 'Owner_State'
```

```
In [160]: df.rename(columns={'City':'Owner_City', 'State':'Owner_State'})
```

Out[160]:

	UniqueID	ParcelID	LandUse	SaleDate	SalePrice	LegalReference	SoldAsVacant	OwnerName
0	2045	007 00 0 125.00	SINGLE FAMILY	2013-04- 09	240000	20130412- 0036474	No	CHRISTOPHER SAMUELSON
1	16918	007 00 0 130.00	SINGLE FAMILY	2014-06- 10	366000	20140619- 0053768	No	CHRISTOPHER SAMUELSON
2	54582	007 00 0 138.00	SINGLE FAMILY	2016-09- 26	435000	20160927- 0101718	No	CHRISTOPHER SAMUELSON
3	43070	007 00 0 143.00	SINGLE FAMILY	2016-01- 29	255000	20160129- 0008913	No	CHRISTOPHER SAMUELSON
4	22714	007 00 0 149.00	SINGLE FAMILY	2014-10- 10	278000	20141015- 0095255	No	CHRISTOPHER SAMUELSON
...
56472	30469	188 10 0A 101.00	SINGLE FAMILY	2015-05- 27	157500	20150608- 0053286	No	CHRISTOPHER SAMUELSON
56473	27707	188 10 0A 107.00	SINGLE FAMILY	2015-03- 02	145000	20150304- 0019013	No	CHRISTOPHER SAMUELSON
56474	52709	188 10 0A 118.00	VACANT RESIDENTIAL LAND	2016-08- 16	234611	20160819- 0087214	Yes	CHRISTOPHER SAMUELSON
56475	54042	188 10 0A 121.00	VACANT RESIDENTIAL LAND	2016-09- 07	93844	20160919- 0098411	Yes	CHRISTOPHER SAMUELSON
56476	54043	188 10 0A 122.00	VACANT RESIDENTIAL LAND	2016-09- 07	93844	20160919- 0098411	Yes	CHRISTOPHER SAMUELSON

56477 rows × 21 columns



SPLITTING UP THE OwnerName Column

```
In [161]: new4=df["OwnerName"].str.split(", ",n=1,expand=True)
```



```
In [162]: df["Owner's L Name"]=new4[0]
```

```
In [163]: df["Owner's F&M Name"]=new4[1]
```

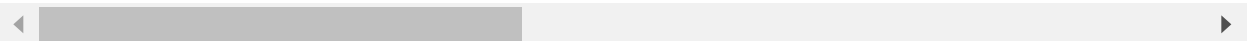
```
In [164]: # Dropping 'OwnerName'
df.drop(columns=["OwnerName"],inplace=True)
```

```
In [165]: df.head(5)
```

Out[165]:

	UniqueID	ParcelID	LandUse	SaleDate	SalePrice	LegalReference	SoldAsVacant	Acreage	Tax
0	2045	007 00 0 125.00	SINGLE FAMILY	2013-04- 09	240000	20130412- 0036474	No	2.3	GE SEF DI
1	16918	007 00 0 130.00	SINGLE FAMILY	2014-06- 10	366000	20140619- 0053768	No	3.5	GE SEF DI
2	54582	007 00 0 138.00	SINGLE FAMILY	2016-09- 26	435000	20160927- 0101718	No	2.9	GE SEF DI
3	43070	007 00 0 143.00	SINGLE FAMILY	2016-01- 29	255000	20160129- 0008913	No	2.6	GE SEF DI
4	22714	007 00 0 149.00	SINGLE FAMILY	2014-10- 10	278000	20141015- 0095255	No	2.0	GE SEF DI

5 rows × 22 columns



```
In [166]: # Listing unique values in SoldAsVacant column
print(df['SoldAsVacant'].unique())
```

```
['No' 'N' 'Yes' 'Y']
```

```
In [167]: # So let's convert 'N' to 'No' and 'Y' to 'Yes'
```

```
In [168]: df[['SoldAsVacant']]=df[['SoldAsVacant']].replace('N','No')
```

```
In [169]: df[['SoldAsVacant']]=df[['SoldAsVacant']].replace('Y','Yes')
```

LET'S RE CHECK UNIQUE VALUES IN SoldAsVacant

```
In [170]: print(df['SoldAsVacant'].unique())  
  
[ 'No'  'Yes' ]
```

YAY....IT WORKED!!!!

There are a lot of things we can do to the data but for the purpose of this project we will end here since what we have done is to introduce some basic methods that can be employed to clean up data. SEE YOU IN MY NEXT PROJECT

Now let us save the final file as a new csv file

```
In [172]: df.to_csv('nash_housing_data_cleaned.csv', index=False)
```

**THE END THE END THE END THE END THE END THE END THE END THE END THE END
THE END THE END THE END THE END THE END THE END**

```
In [ ]:
```